Native bixverse: CPU against GPU, time and memory

Same pipeline and parameters, native on macOS (M1 Max, 8 threads, R 4.6.1). One repeat per run.

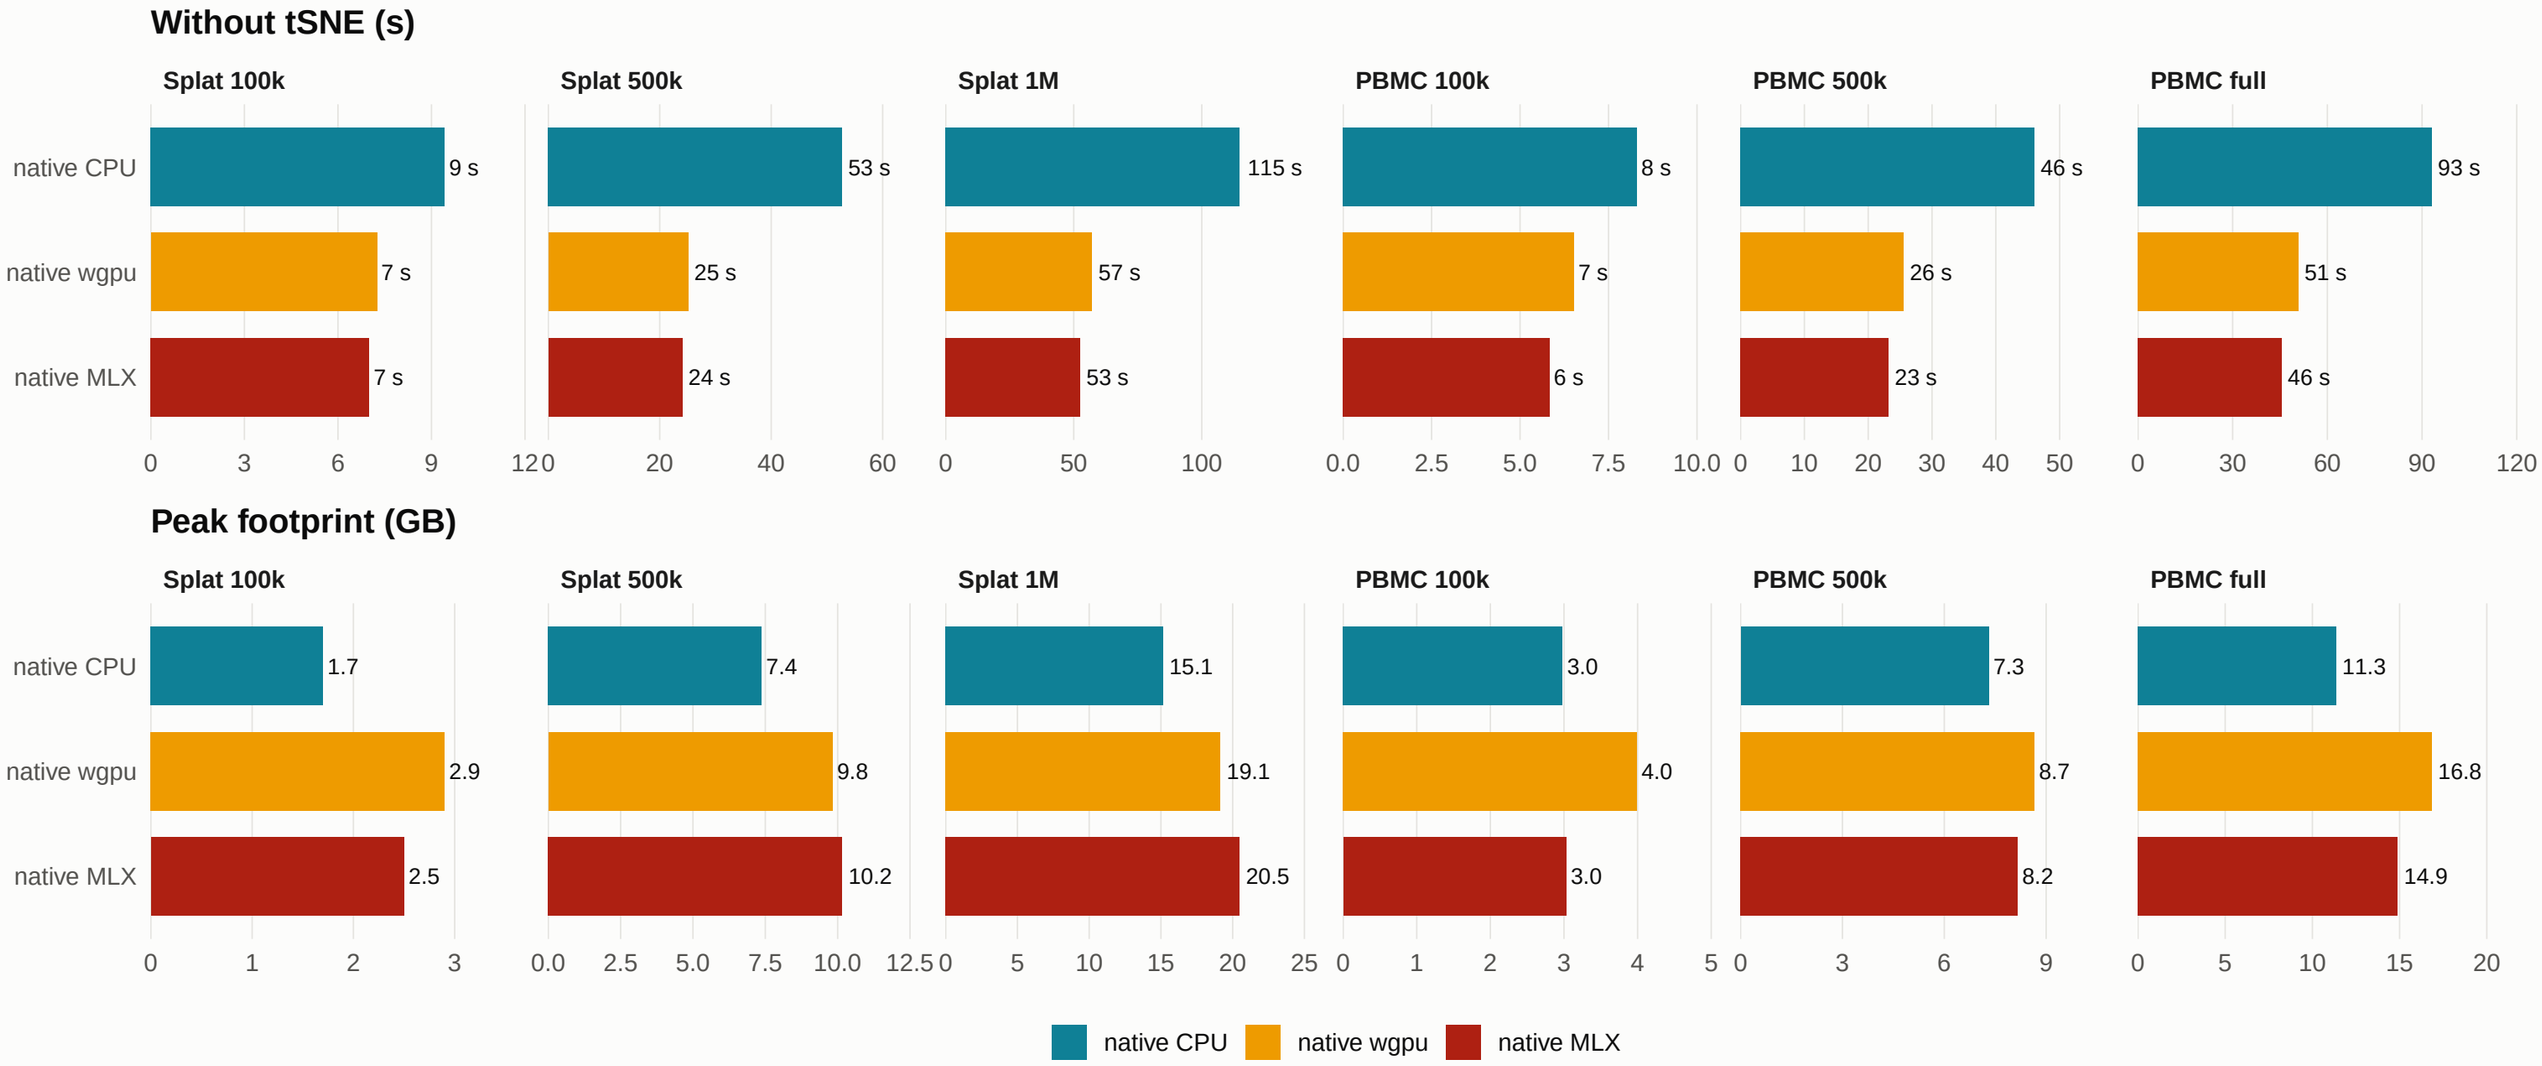

Peak footprint from /usr/bin/time −l over the whole run, tSNE included; no memory cap, so it does not compare with the Docker figures.
At splat 100k the GPU rows' extra memory is CPU heap that grows in the tSNE steps, not GPU buffers; why it is not released is not checked.